

Graph-Level Optimization

Yang-Ge Ma, May 2020

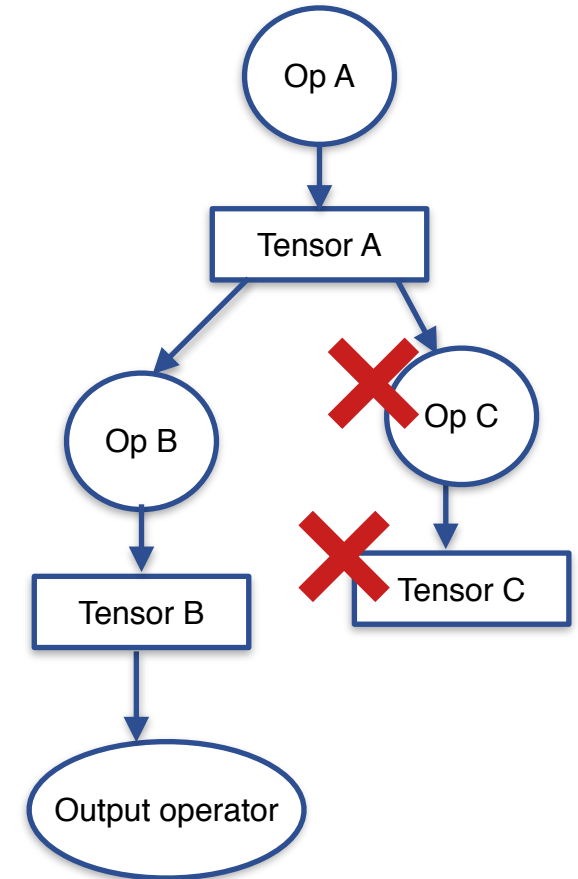


Outline

- Graph-Level Optimizations
 - Target Dependent
 - Target Independent
- Lab
 - Add Optimization Pass
 - Example: EliminateNopTranspose
 - Add Visualization Pass
 - Demo

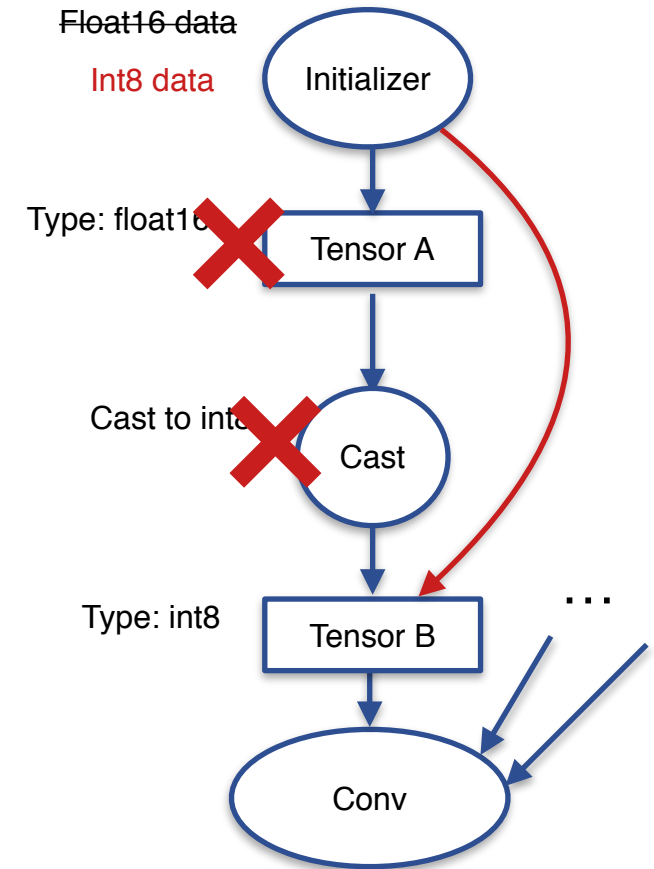
Graph-Level Optimizations

- Manipulation on ONNC IR graphs
- Simplifying IR graphs
 - Node removal
 - **DeadNodeElimination**: Remove nodes whose outputs are not used at all.
 - Layer fusion
 - **EliminateCast**: Fuse Cast operators to preceding initializers.
 - Operator refactoring
 - **ExpandBatchNormalization**: Convert BN to an Add and a Mul.



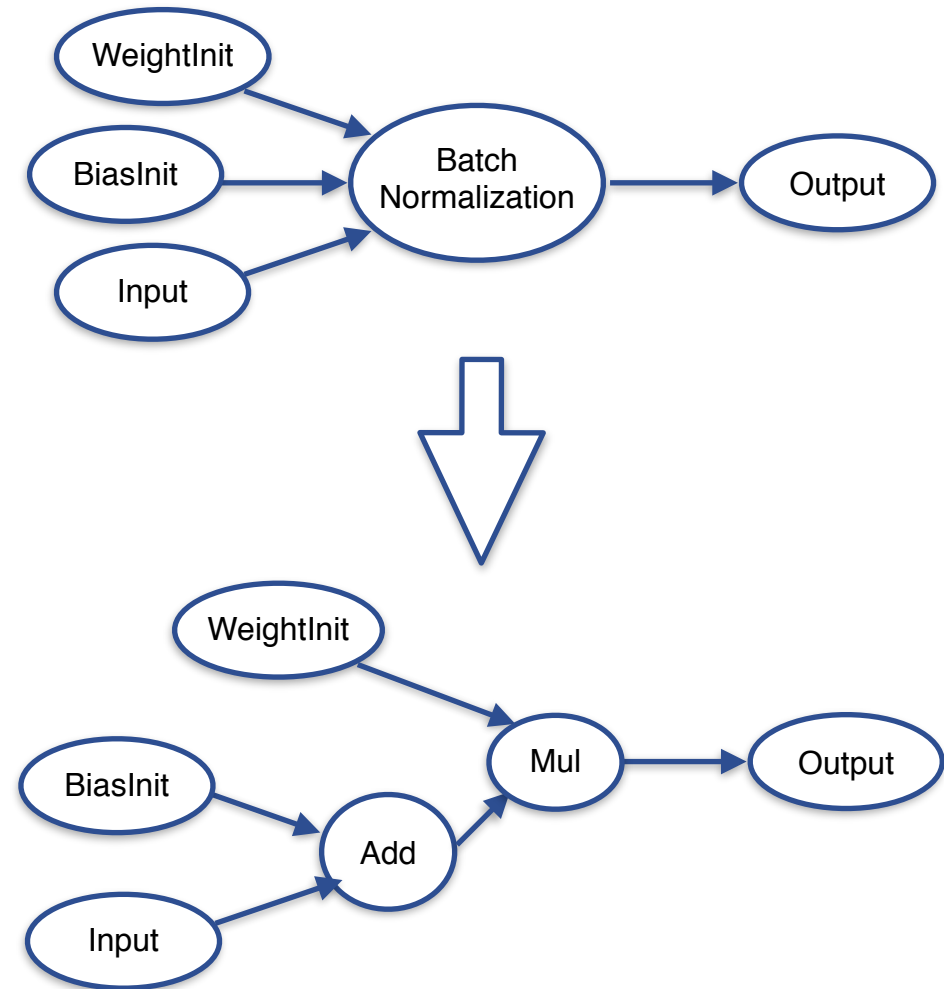
Graph-Level Optimizations

- Manipulation on ONNC IR graphs
- Simplifying IR graphs
 - Node removal
 - **DeadNodeElimination**: Remove nodes whose outputs are not used at all.
 - Layer fusion
 - **EliminateCast**: Fuse Cast operators to preceding initializers.
 - Operator refactoring
 - **ExpandBatchNormalization**: Convert BN to an Add and a Mul.



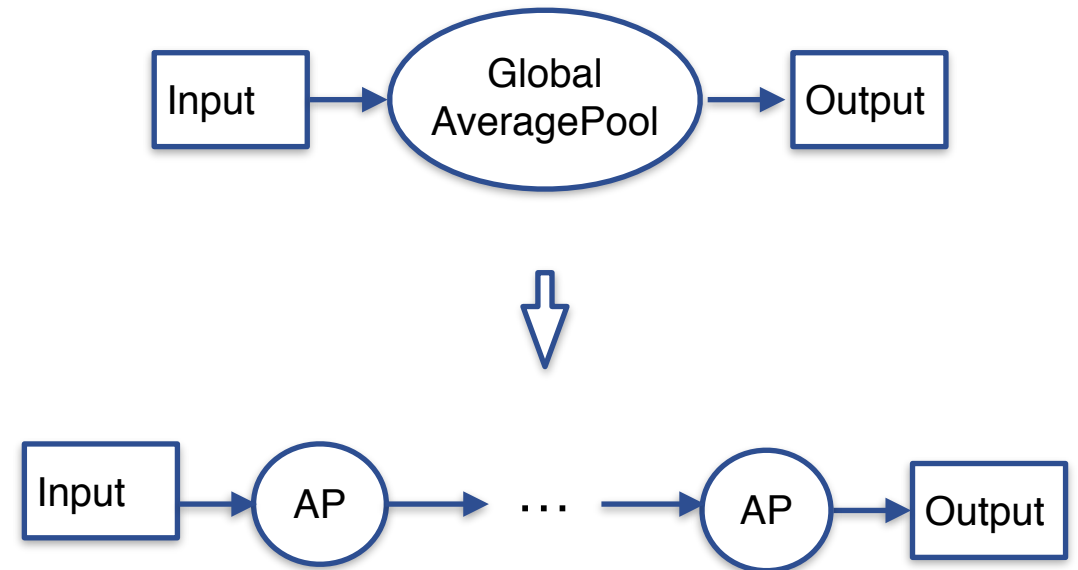
Graph-Level Optimizations

- Manipulation on ONNC IR graphs
- Simplifying IR graphs
 - Node removal
 - **DeadNodeElimination**: Remove nodes whose outputs are not used at all.
 - Layer fusion
 - **EliminateCast**: Fuse Cast operators to preceding initializers.
 - Operator refactoring
 - **ExpandBatchNormalization**: Convert BN to an Add and a Mul.



Graph-Level Optimizations

- Legalization: Convert unsupported operators into supported operators
 - `DivideGlobalAPsIntoAPs`: NVDLA does not support GlobalAveragePool
- Meet hardware constraints
 - `SplitConvPass`: NVDLA's convolution buffer size is limited



Graph-Level Optimizations

- Legalization: Convert unsupported operators into supported operators
 - `DivideGlobalAPsIntoAPs`: NVDLA does not support GlobalAveragePool
- Meet hardware constraints
 - `SplitConvPass`: NVDLA's convolution buffer size is limited

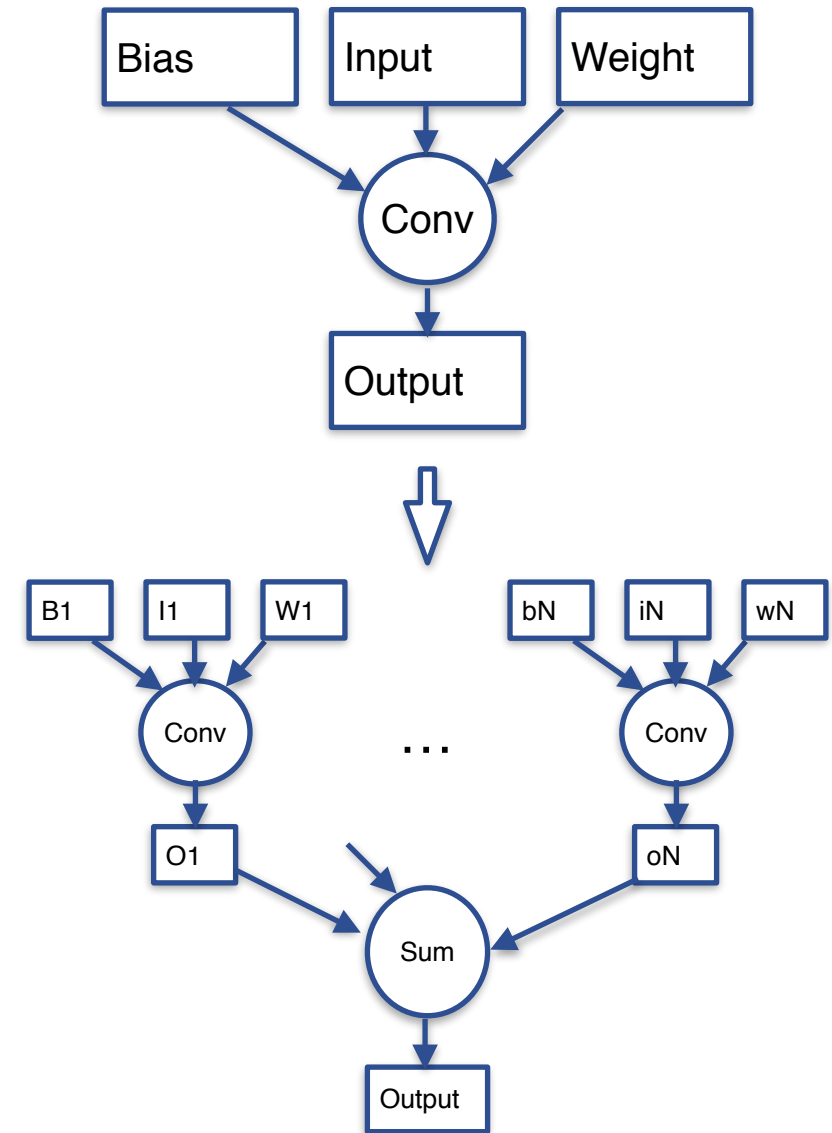



Table of All Currently Supported Optimizations

	Category	Functionality
EliminateIdentity	Node removal	Remove all Identity operators
DeadNodeElimination	Node removal	Remove nodes whose output are totally unused
RemoveTrainingNodes	Node removal	Remove all Dropout nodes yielded during training
EliminateCast	Layer fusion	Fuse Cast operators to preceding initializers
PropagateConstWithDiffShape	Layer fusion	Fuse general reshaping operators to preceding initializers
ExpandBatchNormalization	Operator refactoring	Replace any BatchNormalization operator with an Add and a Mul
ExtractConstToInitializer	Operator refactoring	Make outputs of Const operators initializers
ReplaceGemmByConv	Legalization	Replace Gemm operators, with a Conv and some other operators.
DivideGlobalAPIntoAPs	Legalization	Divide and conquer: break GlobalAPs into AveragePools of smaller windows recursively.
SplitConvPass	Meeting HW constraints	Split input into smaller pieces, do Conv for each of them and accumulate the results.

Lab: Add an Optimization Pass

- EliminateNopTranspose: We demonstrate with an optimization that remove Transpose operator whose permutation is identity.



Workflow Overview

- Steps
 - Implementing a pass
 - Derive `CustomPass` class template
 - Registering the pass to the target backend implementation
 - Add it to the `PassManager`
 - Adding the files of the new pass to the building system

Derive Custom Pass

```
#include <onnc/Core/CustomPass.h>

namespace onnc {
class EliminateNopTranspose: public CustomPass<EliminateNopTranspose>
{
public:
    EliminateNopTranspose() = default;

    // Entry point of the pass
    ReturnTpe runOnModule(Module& pModule) override;
};
} // namespace of onnc
```

Implementation of runOnModule()

```
Pass::ReturnType EliminateNopTranspose::runOnModule(Module& pModule) {  
    using namespace internal;  
    Pass::ReturnType ret = Pass::kModuleNoChanged;  
  
    // Iterate through all ComputeGraphs  
    for (auto it = pModule.cgBegin(); it != pModule.cgEnd(); ++it) {  
        ComputeGraph& cg = *it->value();  
  
        /* Procedure for each compute graph  
         * ...  
         */  
    }  
  
    return ret;  
}
```


Implementation of runOnModule()

```
// In each iteration
```

```
std::vector<ComputeOperator*> removeList;
```

```
for (ComputeOperator& node : cg) {  
    ret = Pass::kModuleChanged;  
    if(isNopTranspose(&node))  
        removeList.emplace_back(&node);  
}
```

```
for (auto* pNopTranspose : removeList) {  
    Value* inV = pNopTranspose->getInput(0);  
    Value* outV = pNopTranspose->getOutput(0);  
    outV->replaceAllUsesWith(*inV);  
    pNopTranspose->removeAllInputs();  
    cg.erase(*pNopTranspose);  
    cg.erase(*outV);  
}
```




Implementation of runOnModule()

// In each iteration

```
std::vector<ComputeOperator*> removeList;

for (ComputeOperator& node : cg) {
    ret = Pass::kModuleChanged;
    if(isNopTranspose(&node))
        removeList.emplace_back(&node);
}

for (auto* pNopTranspose : removeList) {
    Value* inV = pNopTranspose->getInput(0);
    Value* outV = pNopTranspose->getOutput(0);
    outV->replaceAllUsesWith(*inV);
    pNopTranspose->removeAllInputs();
    cg.erase(*pNopTranspose);
    cg.erase(*outV);
}
```




Implementation of runOnModule()

// In each iteration

```
std::vector<ComputeOperator*> removeList;

for (ComputeOperator& node : cg) {
    ret = Pass::kModuleChanged;
    if(isNopTranspose(&node))
        removeList.emplace_back(&node);
}

for (auto* pNopTranspose : removeList) {
    Value* inV = pNopTranspose->getInput(0);
    Value* outV = pNopTranspose->getOutput(0);
    outV->replaceAllUsesWith(*inV);
    pNopTranspose->removeAllInputs();
    cg.erase(*pNopTranspose);
    cg.erase(*outV);
}
```




Implementation of runOnModule()

// In each iteration

```
std::vector<ComputeOperator*> removeList;

for (ComputeOperator& node : cg) {
    ret = Pass::kModuleChanged;
    if(isNopTranspose(&node))
        removeList.emplace_back(&node);
}

for (auto* pNopTranspose : removeList) {
    Value* inV = pNopTranspose->getInput(0);
    Value* outV = pNopTranspose->getOutput(0);
    outV->replaceAllUsesWith(*inV);
    pNopTranspose->removeAllInputs();
    cg.erase(*pNopTranspose);
    cg.erase(*outV);
}
```



Register New Optimization Pass

- In your backend implementation (say NvDlaBackend.cpp):

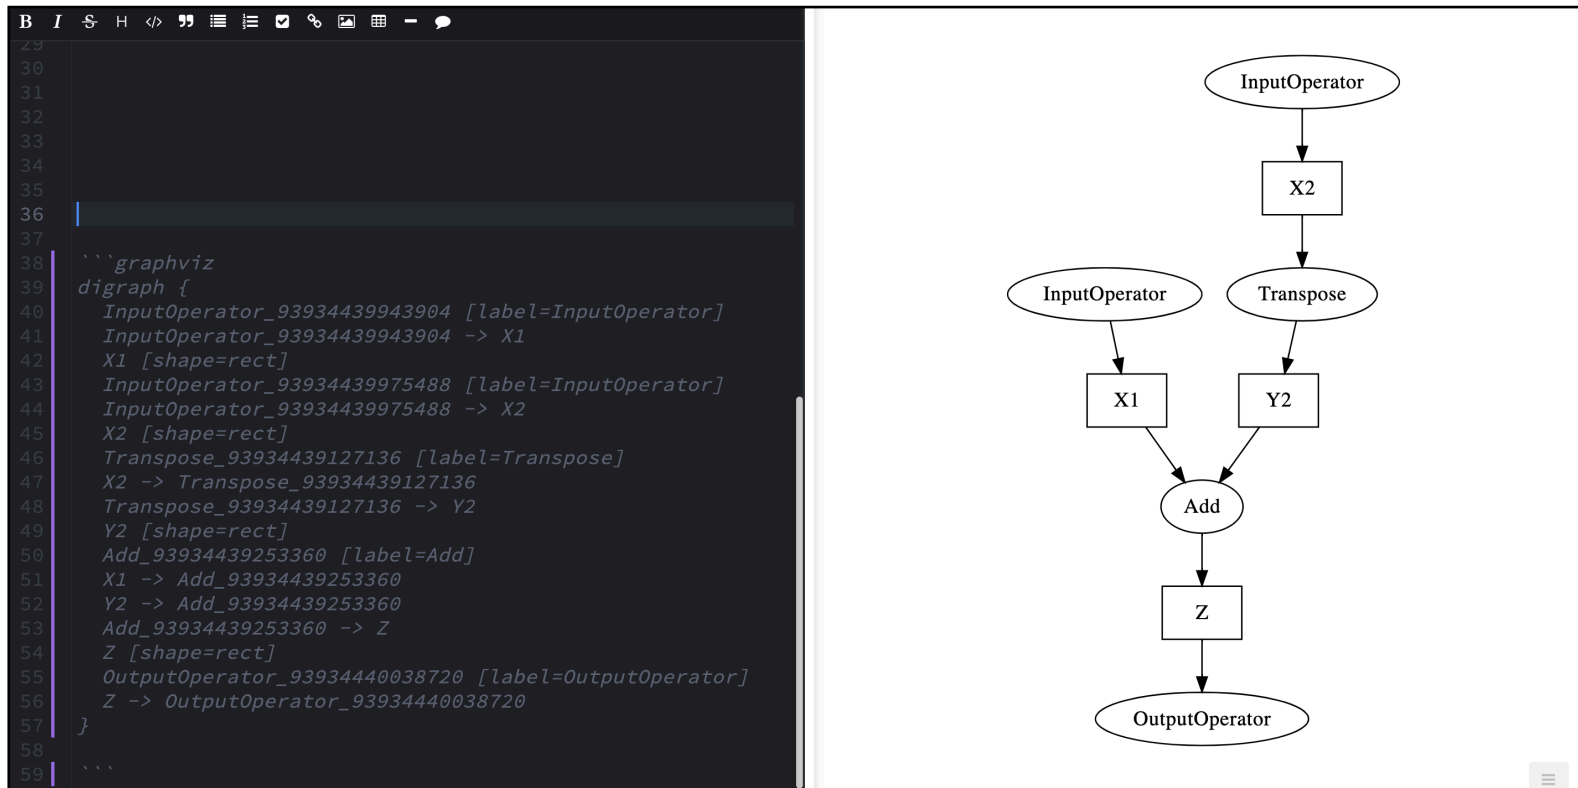
```
+ #include "EliminateNopTranspose.h"
#include <onnc/Transforms/Optimizations/EliminateIdentity.h>
#include <onnc/Transforms/Optimizations/PropagateConstWithDiffShape.h>
NvDlaBackend::addOnncIrOptimization
    (PassManager& passManager, OptimizationOptions& options) {
    // ..
+   passManager.add<EliminateNopTranspose>();
    // ..
}
```

- In CMakeList.txt where you put your Pass implementation:

```
add_libonnc_src(
    GraphvizONNCIRPass.cpp
+   EliminateNopTranspose.cpp
)
```

Lab: Add Visualization Pass

- Graphviz: A graph visualization language supported in markdown
- Implement a pass that dumps the graph in Graphviz format to see the change after optimization.





Skymizer Taiwan Inc.

CONTACT US

E-mail sales@skymizer.com **Tel** +886 2 8797 8337

HQ 12F-2, No.408, Ruiguang Rd., Neihu Dist., Taipei City 11492, Taiwan

BR Center of Innovative Incubator, National Tsing Hua University, Hsinchu Taiwan



skymizer

Boost deep learning accelerator
with compiler technology



<https://skymizer.com>